

# Implementing various MLP Architectures using Keras.

In [1]:

```
import tensorflow as tf
```

In [2]:

```
# checking if the gpu is connected or not
```

```
from tensorflow.python.client import device_lib  
print(device_lib.list_local_devices())
```

```
[name: "/device:CPU:0"  
device_type: "CPU"  
memory_limit: 268435456  
locality {  
}  
incarnation: 9299138210179744872  
, name: "/device:GPU:0"  
device_type: "GPU"  
memory_limit: 4840685568  
locality {  
  bus_id: 1  
  links {  
  }  
}  
incarnation: 2320203755516720165  
physical_device_desc: "device: 0, name: GeForce RTX 2060, pci bus id: 000  
0:01:00.0, compute capability: 7.5"  
]
```

In [3]:

```
from keras.utils import np_utils  
from keras.datasets import mnist  
import seaborn as sns  
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [4]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error-- which we will be
# calling in the later part of the code.
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [5]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [6]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%
d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of test examples :", X_test.shape[0], "and each image is of shape (%d, %d
)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
 Number of test examples : 10000 and each image is of shape (28, 28)

**Now we can see that our input is in the form of a 2-D vector hence we will convert it into a 1-d vector right now that is from 2828 ---> 1784.**

In [7]:

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

## After conversion

In [8]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%
d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d
)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
 Number of training examples : 10000 and each image is of shape (784)

In [9]:

```
# sample data point  
X_train[5]
```

[illegible]

```
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0], dtype=uint8)
```

**Each cell in the above matrix has value between 0 to 255 but as we will be using these cell values for our ML tasks hence lets normalize these values and bring them between 0 and 1**

In [10]:

```
# X => (X - Xmin)/(Xmax-Xmin) = X/255    --- >> note that Xmin = 0

X_train = X_train/255
X_test = X_test/255
```

In [11]:

```
#normalized sample data point  
X_train[5]
```

[illegible]

0. , 0. , 0.30196078, 0.98823529, 0.98823529,  
0.23529412, 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0.10196078, 0.50196078, 0.22745098,  
0.08627451, 0. , 0. , 0. , 0. ,  
0.39215686, 0.98823529, 0.98823529, 0.23529412, 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0.61568627, 0.98823529,  
0.98823529, 0.23529412, 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0.43137255, 0.4745098 , 0.47843137,  
0.4745098 , 0.79215686, 0.98823529, 0.76078431, 0.01176471,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0.03921569, 0.20784314, 0.70196078,  
0.99215686, 0.99215686, 1. , 0.99215686, 0.99215686,  
0.89411765, 0.1372549 , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0.01960784, 0.21176471,  
0.89019608, 0.98823529, 0.95294118, 0.89411765, 0.66666667,  
0.94901961, 0.98823529, 0.98823529, 0.90588235, 0.45882353,  
0.02352941, 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0.02352941, 0.30588235, 0.98823529, 0.98823529, 0.49019608,  
0.23137255, 0. , 0.07058824, 0.81568627, 0.98823529,  
0.98823529, 0.98823529, 0.98823529, 0.34117647, 0.02745098,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0.01960784, 0.52941176, 0.98823529,  
0.98823529, 0.70588235, 0.0627451 , 0. , 0.08235294,  
0.79607843, 0.99215686, 0.96862745, 0.50588235, 0.67843137,  
0.98823529, 0.98823529, 0.72156863, 0.25882353, 0.19215686,  
0.19215686, 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0.01176471,  
0.53333333, 0.98823529, 0.94509804, 0.41568627, 0.06666667,  
0. , 0.20784314, 0.78431373, 0.98823529, 0.84705882,  
0.25490196, 0. , 0.05490196, 0.28235294, 0.63921569,  
0.94509804, 0.98823529, 0.98823529, 0.8745098 , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0.41176471, 0.98823529, 0.94901961,  
0.34509804, 0.07058824, 0.28627451, 0.66666667, 0.95686275,  
0.98823529, 0.49411765, 0.11372549, 0. , 0. ,  
0. , 0. , 0. , 0.34901961, 0.70588235,  
0.70588235, 0.14509804, 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0.90588235, 0.98823529, 0.96078431, 0.80392157, 0.84705882,  
0.98823529, 0.98823529, 0.98823529, 0.48627451, 0.01176471,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0. , 0. ,  
0. , 0. , 0. , 0.81176471, 0.98823529,  
0.98823529, 0.98823529, 0.98823529, 0.69803922, 0.45490196,



[illegible]

**Now let us check the class labels.**

In [12]:

```
# given Label
y_train[5]
```

Out[12]:

2

Each class label will be a number between 0 to 9 hence let us convert them into vector form using one hot encoding ---> ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0] We are doing so because we will be using the categorical cross entropy loss for our model

In [13]:

```
Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)
```

In [14]:

```
#Label after one hot encoding
Y_train[5]
```

Out[14]:

```
array([0., 0., 1., 0., 0., 0., 0., 0., 0., 0.], dtype=float32)
```

In [15]:

```
len(Y_train[5])
```

Out[15]:

```
10
```

## Using Keras to build our Sequential Model

In [16]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [17]:

```
X_train.shape[1]
```

Out[17]:

```
784
```

In [18]:

```
# some model parameters

output_dim = 10 # as our output will be any number from 0 to 9
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

**The model needs to know what input shape it should expect. For this reason, the first layer in a Sequential model (and only the first, because following layers can do automatic shape inference) needs to receive information about its input shape. you can use input\_shape and input\_dim to pass the shape of input.**

## Start building the model

# Model-1 Type -> Models with Relu activation + Adam optimizer + Dropouts + Batch Normalization

## Model - 1 --> 2 Hidden Layers

Model Architecture --> [(input)--(hidden layer 1 (512 neurons))---(hidden layer 2 (128 neurons))---(output layer)]

In [31]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(\theta, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(\theta, \sigma) = N(\theta, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(\theta, \sigma) = N(\theta, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(\theta, \sigma) = N(\theta, 0.120)$ 
```

In [25]:

```
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
```

In [48]:

```

model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_8 (Dense)	(None, 512)	401920
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 128)	65664
batch_normalization_7 (Batch Normalization)	(None, 128)	512
dropout_7 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 10)	1290
=====	=====	=====
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		
=====	=====	=====

In [49]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 3s 51us/step - loss: 0.4832  
- acc: 0.8522 - val\_loss: 0.1638 - val\_acc: 0.9504

Epoch 2/20

60000/60000 [=====] - 2s 40us/step - loss: 0.2508  
- acc: 0.9243 - val\_loss: 0.1246 - val\_acc: 0.9612

Epoch 3/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1981  
- acc: 0.9406 - val\_loss: 0.1068 - val\_acc: 0.9675

Epoch 4/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1788  
- acc: 0.9466 - val\_loss: 0.0953 - val\_acc: 0.9707

Epoch 5/20

60000/60000 [=====] - 2s 41us/step - loss: 0.1536  
- acc: 0.9532 - val\_loss: 0.0893 - val\_acc: 0.9720

Epoch 6/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1432  
- acc: 0.9565 - val\_loss: 0.0788 - val\_acc: 0.9767

Epoch 7/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1322  
- acc: 0.9591 - val\_loss: 0.0802 - val\_acc: 0.9750

Epoch 8/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1228  
- acc: 0.9626 - val\_loss: 0.0739 - val\_acc: 0.9777

Epoch 9/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1097  
- acc: 0.9665 - val\_loss: 0.0748 - val\_acc: 0.9770

Epoch 10/20

60000/60000 [=====] - 2s 40us/step - loss: 0.1106  
- acc: 0.9663 - val\_loss: 0.0696 - val\_acc: 0.9794

Epoch 11/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1029  
- acc: 0.9680 - val\_loss: 0.0728 - val\_acc: 0.9785

Epoch 12/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0974  
- acc: 0.9697 - val\_loss: 0.0670 - val\_acc: 0.9798

Epoch 13/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0931  
- acc: 0.9704 - val\_loss: 0.0646 - val\_acc: 0.9804

Epoch 14/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0843  
- acc: 0.9741 - val\_loss: 0.0688 - val\_acc: 0.9802

Epoch 15/20

60000/60000 [=====] - 3s 43us/step - loss: 0.0848  
- acc: 0.9735 - val\_loss: 0.0680 - val\_acc: 0.9790

Epoch 16/20

60000/60000 [=====] - 2s 39us/step - loss: 0.0809  
- acc: 0.9749 - val\_loss: 0.0638 - val\_acc: 0.9800

Epoch 17/20

60000/60000 [=====] - 2s 39us/step - loss: 0.0785  
- acc: 0.9749 - val\_loss: 0.0633 - val\_acc: 0.9814

Epoch 18/20

60000/60000 [=====] - 3s 44us/step - loss: 0.0747  
- acc: 0.9760 - val\_loss: 0.0585 - val\_acc: 0.9821

Epoch 19/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0730  
- acc: 0.9771 - val\_loss: 0.0599 - val\_acc: 0.9819

Epoch 20/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0672  
- acc: 0.9785 - val\_loss: 0.0621 - val\_acc: 0.9825

In [50]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06209656530671637

Test accuracy: 0.9825



## Let's give a look at the weights which lead to this accuracy of our model.

In [51]:


```
w_after = model_final.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## Model - 1 --> 3 Hidden Layers

**Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))---(output layer)]**

In [52]:

```
model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
batch_normalization_8 (Batch Normalization)	(None, 512)	2048
dropout_8 (Dropout)	(None, 512)	0
dense_12 (Dense)	(None, 128)	65664
batch_normalization_9 (Batch Normalization)	(None, 128)	512
dropout_9 (Dropout)	(None, 128)	0
dense_13 (Dense)	(None, 256)	33024
batch_normalization_10 (Batch Normalization)	(None, 256)	1024
dropout_10 (Dropout)	(None, 256)	0
dense_14 (Dense)	(None, 10)	2570
Total params: 506,762		
Trainable params: 504,970		
Non-trainable params: 1,792		

In [53]:

```
batch_size = 128  
nb_epoch = 40
```

In [54]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/40

60000/60000 [=====] - 4s 69us/step - loss: 0.7553  
- acc: 0.7651 - val\_loss: 0.2128 - val\_acc: 0.9367

Epoch 2/40

60000/60000 [=====] - 3s 52us/step - loss: 0.3450  
- acc: 0.8973 - val\_loss: 0.1563 - val\_acc: 0.9502

Epoch 3/40

60000/60000 [=====] - 3s 54us/step - loss: 0.2711  
- acc: 0.9201 - val\_loss: 0.1286 - val\_acc: 0.9595

Epoch 4/40

60000/60000 [=====] - 3s 52us/step - loss: 0.2319  
- acc: 0.9310 - val\_loss: 0.1067 - val\_acc: 0.9671

Epoch 5/40

60000/60000 [=====] - 3s 51us/step - loss: 0.2034  
- acc: 0.9406 - val\_loss: 0.1019 - val\_acc: 0.9674

Epoch 6/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1789  
- acc: 0.9473 - val\_loss: 0.1002 - val\_acc: 0.9708

Epoch 7/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1682  
- acc: 0.9505 - val\_loss: 0.0938 - val\_acc: 0.9720

Epoch 8/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1603  
- acc: 0.9526 - val\_loss: 0.0859 - val\_acc: 0.9748

Epoch 9/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1463  
- acc: 0.9561 - val\_loss: 0.0824 - val\_acc: 0.9745

Epoch 10/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1376  
- acc: 0.9584 - val\_loss: 0.0818 - val\_acc: 0.9758

Epoch 11/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1288  
- acc: 0.9611 - val\_loss: 0.0812 - val\_acc: 0.9761

Epoch 12/40

60000/60000 [=====] - 3s 53us/step - loss: 0.1222  
- acc: 0.9626 - val\_loss: 0.0798 - val\_acc: 0.9757

Epoch 13/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1189  
- acc: 0.9655 - val\_loss: 0.0784 - val\_acc: 0.9769

Epoch 14/40

60000/60000 [=====] - 3s 51us/step - loss: 0.1096  
- acc: 0.9671 - val\_loss: 0.0717 - val\_acc: 0.9796

Epoch 15/40

60000/60000 [=====] - 3s 49us/step - loss: 0.1064  
- acc: 0.9683 - val\_loss: 0.0737 - val\_acc: 0.9788

Epoch 16/40

60000/60000 [=====] - 3s 49us/step - loss: 0.1036  
- acc: 0.9695 - val\_loss: 0.0768 - val\_acc: 0.9782

Epoch 17/40

60000/60000 [=====] - 3s 51us/step - loss: 0.0972  
- acc: 0.9705 - val\_loss: 0.0716 - val\_acc: 0.9801

Epoch 18/40

60000/60000 [=====] - 3s 52us/step - loss: 0.0933  
- acc: 0.9714 - val\_loss: 0.0723 - val\_acc: 0.9807

Epoch 19/40

60000/60000 [=====] - 3s 51us/step - loss: 0.0927  
- acc: 0.9719 - val\_loss: 0.0687 - val\_acc: 0.9805

Epoch 20/40

60000/60000 [=====] - 3s 51us/step - loss: 0.0882  
- acc: 0.9734 - val\_loss: 0.0685 - val\_acc: 0.9812

```
Epoch 21/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0828
- acc: 0.9745 - val_loss: 0.0657 - val_acc: 0.9821
Epoch 22/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0822
- acc: 0.9749 - val_loss: 0.0692 - val_acc: 0.9811
Epoch 23/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0827
- acc: 0.9755 - val_loss: 0.0666 - val_acc: 0.9812
Epoch 24/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0763
- acc: 0.9766 - val_loss: 0.0654 - val_acc: 0.9809
Epoch 25/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0734
- acc: 0.9775 - val_loss: 0.0654 - val_acc: 0.9826
Epoch 26/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0710
- acc: 0.9781 - val_loss: 0.0646 - val_acc: 0.9823
Epoch 27/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0696
- acc: 0.9795 - val_loss: 0.0624 - val_acc: 0.9812
Epoch 28/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0702
- acc: 0.9786 - val_loss: 0.0680 - val_acc: 0.9815
Epoch 29/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0705
- acc: 0.9789 - val_loss: 0.0628 - val_acc: 0.9827
Epoch 30/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0675
- acc: 0.9797 - val_loss: 0.0643 - val_acc: 0.9832
Epoch 31/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0639
- acc: 0.9804 - val_loss: 0.0650 - val_acc: 0.9831
Epoch 32/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0634
- acc: 0.9800 - val_loss: 0.0642 - val_acc: 0.9828
Epoch 33/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0632
- acc: 0.9806 - val_loss: 0.0658 - val_acc: 0.9825
Epoch 34/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0616
- acc: 0.9807 - val_loss: 0.0623 - val_acc: 0.9831
Epoch 35/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0553
- acc: 0.9830 - val_loss: 0.0643 - val_acc: 0.9828
Epoch 36/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0584
- acc: 0.9825 - val_loss: 0.0644 - val_acc: 0.9830
Epoch 37/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0557
- acc: 0.9828 - val_loss: 0.0601 - val_acc: 0.9843
Epoch 38/40
60000/60000 [=====] - 3s 51us/step - loss: 0.0568
- acc: 0.9821 - val_loss: 0.0559 - val_acc: 0.9847
Epoch 39/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0531
- acc: 0.9832 - val_loss: 0.0583 - val_acc: 0.9849
Epoch 40/40
60000/60000 [=====] - 3s 52us/step - loss: 0.0520
- acc: 0.9839 - val_loss: 0.0645 - val_acc: 0.9833
```

In [55]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# Loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06454481310489937

Test accuracy: 0.9833



**Let's give a look at the weights which lead to this accuracy of our model.**

In [56]:

```
w_after = model_final.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()

plt.subplot(1, 4, 1)
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.title("Weight matrices after training")
plt.subplot(1, 4, 2)
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 4, 3)
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')

plt.show()
```



Weight matrices after training



## Model - 1 --> 5 Hidden Layers

**Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))--(hidden layer 4 (1024 neurons))----(hidden layer 5 (2048 neurons))---(output layer)]**

In [57]:

```
model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(1024, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(2048, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())
model_final.add(Dropout(0.5))

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()
```

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 512)	401920
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
dense_16 (Dense)	(None, 128)	65664
batch_normalization_12 (Batch Normalization)	(None, 128)	512
dropout_12 (Dropout)	(None, 128)	0
dense_17 (Dense)	(None, 256)	33024
batch_normalization_13 (Batch Normalization)	(None, 256)	1024
dropout_13 (Dropout)	(None, 256)	0
dense_18 (Dense)	(None, 1024)	263168
batch_normalization_14 (Batch Normalization)	(None, 1024)	4096
dropout_14 (Dropout)	(None, 1024)	0
dense_19 (Dense)	(None, 2048)	2099200
batch_normalization_15 (Batch Normalization)	(None, 2048)	8192
dropout_15 (Dropout)	(None, 2048)	0
dense_20 (Dense)	(None, 10)	20490
Total params: 2,899,338		
Trainable params: 2,891,402		
Non-trainable params: 7,936		

In [58]:

```
batch_size = 256
nb_epoch = 60
```

In [59]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/60

60000/60000 [=====] - 5s 76us/step - loss: 1.3446  
- acc: 0.5888 - val\_loss: 0.4966 - val\_acc: 0.8424

Epoch 2/60

60000/60000 [=====] - 3s 48us/step - loss: 0.5576  
- acc: 0.8277 - val\_loss: 0.2985 - val\_acc: 0.9067

Epoch 3/60

60000/60000 [=====] - 3s 47us/step - loss: 0.4277  
- acc: 0.8710 - val\_loss: 0.2539 - val\_acc: 0.9254

Epoch 4/60

60000/60000 [=====] - 3s 47us/step - loss: 0.3590  
- acc: 0.8942 - val\_loss: 0.2306 - val\_acc: 0.9334

Epoch 5/60

60000/60000 [=====] - 3s 47us/step - loss: 0.3114  
- acc: 0.9078 - val\_loss: 0.1989 - val\_acc: 0.9455

Epoch 6/60

60000/60000 [=====] - 3s 49us/step - loss: 0.2770  
- acc: 0.9191 - val\_loss: 0.1746 - val\_acc: 0.9520

Epoch 7/60

60000/60000 [=====] - 3s 47us/step - loss: 0.2518  
- acc: 0.9277 - val\_loss: 0.1642 - val\_acc: 0.9563

Epoch 8/60

60000/60000 [=====] - 3s 47us/step - loss: 0.2302  
- acc: 0.9326 - val\_loss: 0.1355 - val\_acc: 0.9624

Epoch 9/60

60000/60000 [=====] - 3s 46us/step - loss: 0.2146  
- acc: 0.9383 - val\_loss: 0.1395 - val\_acc: 0.9619

Epoch 10/60

60000/60000 [=====] - 3s 47us/step - loss: 0.1981  
- acc: 0.9425 - val\_loss: 0.1298 - val\_acc: 0.9642

Epoch 11/60

60000/60000 [=====] - 3s 46us/step - loss: 0.1894  
- acc: 0.9453 - val\_loss: 0.1244 - val\_acc: 0.9677

Epoch 12/60

60000/60000 [=====] - 3s 46us/step - loss: 0.1773  
- acc: 0.9484 - val\_loss: 0.1222 - val\_acc: 0.9696

Epoch 13/60

60000/60000 [=====] - 3s 47us/step - loss: 0.1686  
- acc: 0.9519 - val\_loss: 0.1227 - val\_acc: 0.9693

Epoch 14/60

60000/60000 [=====] - 3s 46us/step - loss: 0.1620  
- acc: 0.9530 - val\_loss: 0.1176 - val\_acc: 0.9706

Epoch 15/60

60000/60000 [=====] - 3s 47us/step - loss: 0.1559  
- acc: 0.9550 - val\_loss: 0.1112 - val\_acc: 0.9726

Epoch 16/60

60000/60000 [=====] - 3s 47us/step - loss: 0.1481  
- acc: 0.9572 - val\_loss: 0.1167 - val\_acc: 0.9700

Epoch 17/60

60000/60000 [=====] - 3s 47us/step - loss: 0.1424  
- acc: 0.9594 - val\_loss: 0.0982 - val\_acc: 0.9758

Epoch 18/60

60000/60000 [=====] - 3s 46us/step - loss: 0.1356  
- acc: 0.9608 - val\_loss: 0.0977 - val\_acc: 0.9752

Epoch 19/60

60000/60000 [=====] - 3s 46us/step - loss: 0.1335  
- acc: 0.9608 - val\_loss: 0.0914 - val\_acc: 0.9766

Epoch 20/60

60000/60000 [=====] - 3s 46us/step - loss: 0.1278  
- acc: 0.9639 - val\_loss: 0.0896 - val\_acc: 0.9786

Epoch 21/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.1266  
- acc: 0.9641 - val\_loss: 0.0959 - val\_acc: 0.9762

Epoch 22/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.1165  
- acc: 0.9665 - val\_loss: 0.0937 - val\_acc: 0.9780

Epoch 23/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.1137  
- acc: 0.9669 - val\_loss: 0.0951 - val\_acc: 0.9780

Epoch 24/60  
60000/60000 [=====] - 3s 47us/step - loss: 0.1101  
- acc: 0.9682 - val\_loss: 0.0877 - val\_acc: 0.9797

Epoch 25/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.1106  
- acc: 0.9676 - val\_loss: 0.0937 - val\_acc: 0.9779

Epoch 26/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.1135  
- acc: 0.9680 - val\_loss: 0.0900 - val\_acc: 0.9794

Epoch 27/60  
60000/60000 [=====] - 3s 47us/step - loss: 0.1032  
- acc: 0.9701 - val\_loss: 0.0982 - val\_acc: 0.9782

Epoch 28/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.1015  
- acc: 0.9705 - val\_loss: 0.0950 - val\_acc: 0.9777

Epoch 29/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0981  
- acc: 0.9724 - val\_loss: 0.0938 - val\_acc: 0.9780

Epoch 30/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0940  
- acc: 0.9732 - val\_loss: 0.0894 - val\_acc: 0.9798

Epoch 31/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0869  
- acc: 0.9750 - val\_loss: 0.0885 - val\_acc: 0.9800

Epoch 32/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0901  
- acc: 0.9738 - val\_loss: 0.0859 - val\_acc: 0.9805

Epoch 33/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0853  
- acc: 0.9754 - val\_loss: 0.0921 - val\_acc: 0.9793

Epoch 34/60  
60000/60000 [=====] - 3s 47us/step - loss: 0.0827  
- acc: 0.9757 - val\_loss: 0.0850 - val\_acc: 0.9811

Epoch 35/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0820  
- acc: 0.9757 - val\_loss: 0.0839 - val\_acc: 0.9800

Epoch 36/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0823  
- acc: 0.9767 - val\_loss: 0.0811 - val\_acc: 0.9817

Epoch 37/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0772  
- acc: 0.9782 - val\_loss: 0.0789 - val\_acc: 0.9820

Epoch 38/60  
60000/60000 [=====] - 3s 47us/step - loss: 0.0772  
- acc: 0.9776 - val\_loss: 0.0810 - val\_acc: 0.9815

Epoch 39/60  
60000/60000 [=====] - 3s 47us/step - loss: 0.0748  
- acc: 0.9781 - val\_loss: 0.0790 - val\_acc: 0.9817

Epoch 40/60  
60000/60000 [=====] - 3s 46us/step - loss: 0.0694  
- acc: 0.9794 - val\_loss: 0.0846 - val\_acc: 0.9812

Epoch 41/60

```
60000/60000 [=====] - 3s 46us/step - loss: 0.0711
- acc: 0.9790 - val_loss: 0.0765 - val_acc: 0.9824
Epoch 42/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0694
- acc: 0.9796 - val_loss: 0.0883 - val_acc: 0.9816
Epoch 43/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0703
- acc: 0.9795 - val_loss: 0.0839 - val_acc: 0.9819
Epoch 44/60
60000/60000 [=====] - 3s 47us/step - loss: 0.0689
- acc: 0.9796 - val_loss: 0.0771 - val_acc: 0.9821
Epoch 45/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0645
- acc: 0.9812 - val_loss: 0.0761 - val_acc: 0.9829
Epoch 46/60
60000/60000 [=====] - 3s 47us/step - loss: 0.0671
- acc: 0.9798 - val_loss: 0.0816 - val_acc: 0.9819
Epoch 47/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0625
- acc: 0.9813 - val_loss: 0.0786 - val_acc: 0.9826
Epoch 48/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0590
- acc: 0.9825 - val_loss: 0.0838 - val_acc: 0.9817
Epoch 49/60
60000/60000 [=====] - 3s 47us/step - loss: 0.0579
- acc: 0.9831 - val_loss: 0.0833 - val_acc: 0.9807
Epoch 50/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0604
- acc: 0.9825 - val_loss: 0.0851 - val_acc: 0.9814
Epoch 51/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0572
- acc: 0.9833 - val_loss: 0.0809 - val_acc: 0.9834
Epoch 52/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0563
- acc: 0.9829 - val_loss: 0.0815 - val_acc: 0.9834
Epoch 53/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0561
- acc: 0.9834 - val_loss: 0.0802 - val_acc: 0.9835
Epoch 54/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0575
- acc: 0.9831 - val_loss: 0.0796 - val_acc: 0.9830
Epoch 55/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0536
- acc: 0.9843 - val_loss: 0.0735 - val_acc: 0.9838
Epoch 56/60
60000/60000 [=====] - 3s 47us/step - loss: 0.0527
- acc: 0.9844 - val_loss: 0.0745 - val_acc: 0.9847
Epoch 57/60
60000/60000 [=====] - 3s 46us/step - loss: 0.0542
- acc: 0.9842 - val_loss: 0.0787 - val_acc: 0.9832
Epoch 58/60
60000/60000 [=====] - 3s 47us/step - loss: 0.0528
- acc: 0.9850 - val_loss: 0.0776 - val_acc: 0.9835
Epoch 59/60
60000/60000 [=====] - 3s 47us/step - loss: 0.0491
- acc: 0.9853 - val_loss: 0.0804 - val_acc: 0.9843
Epoch 60/60
60000/60000 [=====] - 3s 48us/step - loss: 0.0502
- acc: 0.9852 - val_loss: 0.0763 - val_acc: 0.9841
```

In [60]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07631430160264717

Test accuracy: 0.9841





**Let's give a look at the weights which lead to this accuracy of our model.**

In [61]:

```
w_after = model_final.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()

plt.subplot(1, 6, 1)
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('HL 1')

plt.title("Weight matrices after training")
plt.subplot(1, 6, 2)
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('HL 2')

plt.subplot(1, 6, 3)
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('HL 3 ')


plt.subplot(1, 6, 4)
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('HL 4 ')

plt.subplot(1, 6, 5)
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('HL 5 ')

plt.subplot(1, 6, 6)
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')

plt.show()
```

Weight matrices after training



**Now Randomly trying different MLP architectures with different configurations.**

**RELU ACTIVATION + ADAM OPTIMIZER + Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))--(hidden layer 4 (1024 neurons))----(hidden layer 5 (2048 neurons))---(output layer)]**

In [62]:

```

model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model_final.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model_final.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model_final.add(Dense(1024, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model_final.add(Dense(2048, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()

```

Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 512)	401920
dense_22 (Dense)	(None, 128)	65664
dense_23 (Dense)	(None, 256)	33024
dense_24 (Dense)	(None, 1024)	263168
dense_25 (Dense)	(None, 2048)	2099200
dense_26 (Dense)	(None, 10)	20490
Total params: 2,883,466		
Trainable params: 2,883,466		
Non-trainable params: 0		

In [63]:

```
nb_epoch = 40
```

In [64]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/40

60000/60000 [=====] - 3s 44us/step - loss: 14.534  
2 - acc: 0.0983 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 2/40

60000/60000 [=====] - 2s 27us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 3/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 4/40

60000/60000 [=====] - 2s 29us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 5/40

60000/60000 [=====] - 2s 29us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 6/40

60000/60000 [=====] - 2s 29us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 7/40

60000/60000 [=====] - 2s 29us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 8/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 9/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 10/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 11/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 12/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 13/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 14/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 15/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 16/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 17/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 18/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 19/40

60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 20/40

60000/60000 [=====] - 2s 29us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 21/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 22/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 23/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 24/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 25/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 26/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 27/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 28/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 29/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 30/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 31/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 32/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 33/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 34/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 35/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 36/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 37/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 38/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 39/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

Epoch 40/40  
60000/60000 [=====] - 2s 28us/step - loss: 14.528  
3 - acc: 0.0986 - val\_loss: 14.5740 - val\_acc: 0.0958

In [65]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

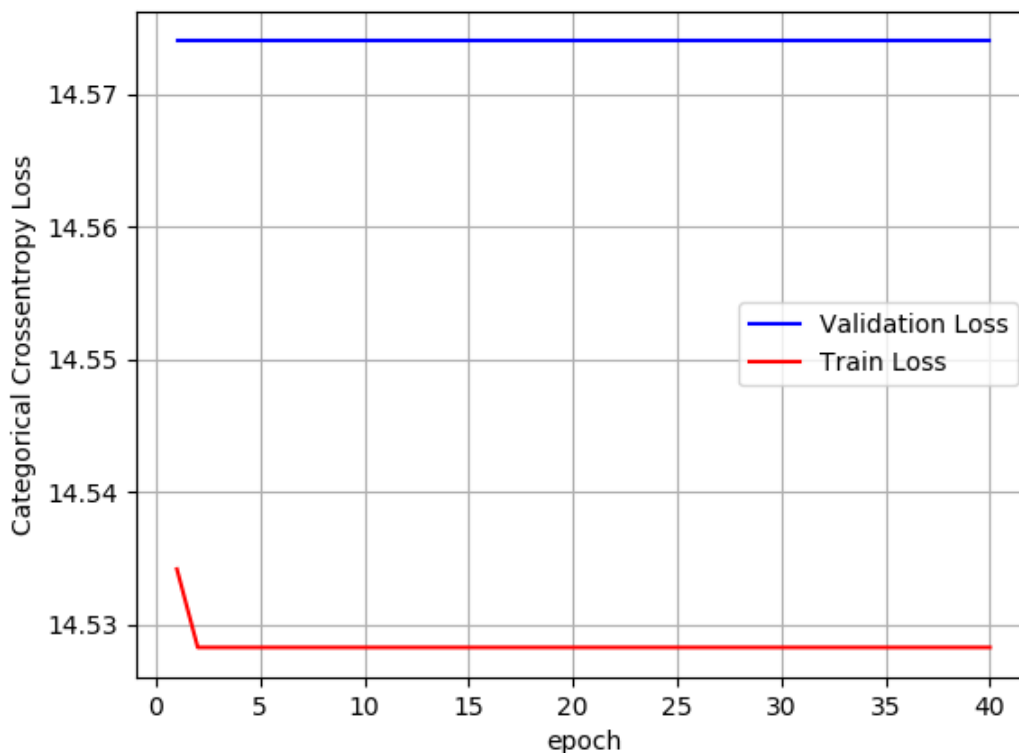
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 14.573981648254394

Test accuracy: 0.0958





**We can clearly see that as we removed the drop out layers and batch normalization we noticed a very poor accuracy for same number of epochs.**

**RELU ACTIVATION + SGD OPTIMIZER + Batch Normalization + Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))--(hidden layer 4 (1024 neurons))----(hidden layer 5 (2048 neurons))---(output layer)]**

In [67]:

```

model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_final.add(BatchNormalization())

model_final.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())

model_final.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())

model_final.add(Dense(1024, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())

model_final.add(Dense(2048, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()

```

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 512)	401920
batch_normalization_16 (Batch Normalization)	(None, 512)	2048
dense_29 (Dense)	(None, 128)	65664
batch_normalization_17 (Batch Normalization)	(None, 128)	512
dense_30 (Dense)	(None, 256)	33024
batch_normalization_18 (Batch Normalization)	(None, 256)	1024
dense_31 (Dense)	(None, 1024)	263168
batch_normalization_19 (Batch Normalization)	(None, 1024)	4096
dense_32 (Dense)	(None, 2048)	2099200
batch_normalization_20 (Batch Normalization)	(None, 2048)	8192
dense_33 (Dense)	(None, 10)	20490
Total params: 2,899,338		
Trainable params: 2,891,402		
Non-trainable params: 7,936		

In [68]:

```
nb_epoch = 40  
batch_size = 128
```

In [69]:

```
model_final.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/40

60000/60000 [=====] - 5s 90us/step - loss: 0.4617  
- acc: 0.8571 - val\_loss: 0.2238 - val\_acc: 0.9308

Epoch 2/40

60000/60000 [=====] - 4s 68us/step - loss: 0.1828  
- acc: 0.9455 - val\_loss: 0.1717 - val\_acc: 0.9482

Epoch 3/40

60000/60000 [=====] - 4s 66us/step - loss: 0.1333  
- acc: 0.9603 - val\_loss: 0.1496 - val\_acc: 0.9549

Epoch 4/40

60000/60000 [=====] - 4s 66us/step - loss: 0.1038  
- acc: 0.9696 - val\_loss: 0.1361 - val\_acc: 0.9595

Epoch 5/40

60000/60000 [=====] - 4s 67us/step - loss: 0.0853  
- acc: 0.9755 - val\_loss: 0.1279 - val\_acc: 0.9608

Epoch 6/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0699  
- acc: 0.9804 - val\_loss: 0.1229 - val\_acc: 0.9628

Epoch 7/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0603  
- acc: 0.9825 - val\_loss: 0.1178 - val\_acc: 0.9646

Epoch 8/40

60000/60000 [=====] - 4s 66us/step - loss: 0.0509  
- acc: 0.9857 - val\_loss: 0.1157 - val\_acc: 0.9658

Epoch 9/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0439  
- acc: 0.9884 - val\_loss: 0.1119 - val\_acc: 0.9673

Epoch 10/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0391  
- acc: 0.9898 - val\_loss: 0.1093 - val\_acc: 0.9681

Epoch 11/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0343  
- acc: 0.9917 - val\_loss: 0.1079 - val\_acc: 0.9689

Epoch 12/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0304  
- acc: 0.9925 - val\_loss: 0.1078 - val\_acc: 0.9680

Epoch 13/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0269  
- acc: 0.9936 - val\_loss: 0.1051 - val\_acc: 0.9686

Epoch 14/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0247  
- acc: 0.9945 - val\_loss: 0.1054 - val\_acc: 0.9691

Epoch 15/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0223  
- acc: 0.9951 - val\_loss: 0.1035 - val\_acc: 0.9700

Epoch 16/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0198  
- acc: 0.9961 - val\_loss: 0.1030 - val\_acc: 0.9700

Epoch 17/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0176  
- acc: 0.9969 - val\_loss: 0.1040 - val\_acc: 0.9702

Epoch 18/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0162  
- acc: 0.9971 - val\_loss: 0.1027 - val\_acc: 0.9703

Epoch 19/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0151  
- acc: 0.9975 - val\_loss: 0.1027 - val\_acc: 0.9709

Epoch 20/40

60000/60000 [=====] - 4s 65us/step - loss: 0.0142  
- acc: 0.9974 - val\_loss: 0.1012 - val\_acc: 0.9720

Epoch 21/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0125  
- acc: 0.9981 - val\_loss: 0.1011 - val\_acc: 0.9715

Epoch 22/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0122  
- acc: 0.9978 - val\_loss: 0.1039 - val\_acc: 0.9710

Epoch 23/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0111  
- acc: 0.9983 - val\_loss: 0.1025 - val\_acc: 0.9714

Epoch 24/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0103  
- acc: 0.9986 - val\_loss: 0.1020 - val\_acc: 0.9719

Epoch 25/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0105  
- acc: 0.9983 - val\_loss: 0.1024 - val\_acc: 0.9710

Epoch 26/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0089  
- acc: 0.9989 - val\_loss: 0.1024 - val\_acc: 0.9718

Epoch 27/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0083  
- acc: 0.9990 - val\_loss: 0.1025 - val\_acc: 0.9709

Epoch 28/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0081  
- acc: 0.9988 - val\_loss: 0.1017 - val\_acc: 0.9714

Epoch 29/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0079  
- acc: 0.9990 - val\_loss: 0.1019 - val\_acc: 0.9713

Epoch 30/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0074  
- acc: 0.9992 - val\_loss: 0.1034 - val\_acc: 0.9707

Epoch 31/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0071  
- acc: 0.9992 - val\_loss: 0.1021 - val\_acc: 0.9717

Epoch 32/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0068  
- acc: 0.9992 - val\_loss: 0.1027 - val\_acc: 0.9719

Epoch 33/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0064  
- acc: 0.9993 - val\_loss: 0.1024 - val\_acc: 0.9720

Epoch 34/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0056  
- acc: 0.9995 - val\_loss: 0.1032 - val\_acc: 0.9713

Epoch 35/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0058  
- acc: 0.9994 - val\_loss: 0.1036 - val\_acc: 0.9716

Epoch 36/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0056  
- acc: 0.9993 - val\_loss: 0.1032 - val\_acc: 0.9725

Epoch 37/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0054  
- acc: 0.9993 - val\_loss: 0.1028 - val\_acc: 0.9719

Epoch 38/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0052  
- acc: 0.9994 - val\_loss: 0.1046 - val\_acc: 0.9718

Epoch 39/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0049  
- acc: 0.9995 - val\_loss: 0.1029 - val\_acc: 0.9724

Epoch 40/40  
60000/60000 [=====] - 4s 65us/step - loss: 0.0050  
- acc: 0.9996 - val\_loss: 0.1020 - val\_acc: 0.9729

In [70]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10197011248232156

Test accuracy: 0.9729



**Sigmoid activation +adam optimizer+ Dropout(with rate = 0.75) +  
Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--  
(hidden layer 2(128 neurons))---(hidden layer 3 (256 neurons))---  
(output layer)]**



In [34]:

```
model_final = Sequential()
model_final.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_final.add(Dropout(0.75))

model_final.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(Dropout(0.75))

model_final.add(Dense(256, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(Dropout(0.75))

model_final.add(Dense(output_dim, activation='softmax'))
model_final.summary()
```

WARNING: Logging before flag parsing goes to stderr.

W1002 10:26:23.843827 9664 deprecation\_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:74: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

W1002 10:26:23.882996 9664 deprecation\_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

W1002 10:26:23.893732 9664 deprecation\_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:4115: The name tf.random\_normal is deprecated. Please use tf.random.normal instead.

W1002 10:26:23.910495 9664 deprecation\_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:133: The name tf.placeholder\_with\_default is deprecated. Please use tf.compat.v1.placeholder\_with\_default instead.

W1002 10:26:23.916351 9664 deprecation.py:506] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

W1002 10:26:23.917327 9664 nn\_ops.py:4224] Large dropout rate: 0.75 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep\_prob. Please ensure that this is intended.

W1002 10:26:23.939802 9664 nn\_ops.py:4224] Large dropout rate: 0.75 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep\_prob. Please ensure that this is intended.

W1002 10:26:23.959618 9664 nn\_ops.py:4224] Large dropout rate: 0.75 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep\_prob. Please ensure that this is intended.

W1002 10:26:23.968427 9664 deprecation\_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:4138: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_1 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 128)	65664
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 256)	33024
dropout_3 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570
=====	=====	=====
Total params: 503,178		
Trainable params: 503,178		
Non-trainable params: 0		
=====		

In [35]:

```
nb_epoch = 40
```

In [36]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_final.fit(X_train, Y_train, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
W1002 10:28:01.442343 9664 deprecation_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
W1002 10:28:01.545196 9664 deprecation.py:323] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
```

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/40

60000/60000 [=====] - 8s 128us/step - loss: 2.428  
7 - acc: 0.1483 - val\_loss: 1.4073 - val\_acc: 0.5483

Epoch 2/40

60000/60000 [=====] - 6s 99us/step - loss: 1.5549  
- acc: 0.4082 - val\_loss: 0.8020 - val\_acc: 0.7208

Epoch 3/40

60000/60000 [=====] - 6s 99us/step - loss: 1.1108  
- acc: 0.5945 - val\_loss: 0.5479 - val\_acc: 0.8673

Epoch 4/40

60000/60000 [=====] - 6s 101us/step - loss: 0.894  
7 - acc: 0.6925 - val\_loss: 0.4284 - val\_acc: 0.8903

Epoch 5/40

60000/60000 [=====] - 6s 99us/step - loss: 0.7512  
- acc: 0.7584 - val\_loss: 0.3424 - val\_acc: 0.9213

Epoch 6/40

60000/60000 [=====] - 6s 100us/step - loss: 0.652  
9 - acc: 0.8037 - val\_loss: 0.2887 - val\_acc: 0.9314

Epoch 7/40

60000/60000 [=====] - 6s 99us/step - loss: 0.5830  
- acc: 0.8317 - val\_loss: 0.2646 - val\_acc: 0.9368

Epoch 8/40

60000/60000 [=====] - 6s 97us/step - loss: 0.5192  
- acc: 0.8549 - val\_loss: 0.2496 - val\_acc: 0.9413

Epoch 9/40

60000/60000 [=====] - 6s 97us/step - loss: 0.4765  
- acc: 0.8712 - val\_loss: 0.2278 - val\_acc: 0.9447

Epoch 10/40

60000/60000 [=====] - 6s 99us/step - loss: 0.4411  
- acc: 0.8819 - val\_loss: 0.2166 - val\_acc: 0.9492

Epoch 11/40

60000/60000 [=====] - 6s 99us/step - loss: 0.4183  
- acc: 0.8907 - val\_loss: 0.1989 - val\_acc: 0.9534

Epoch 12/40

60000/60000 [=====] - 6s 99us/step - loss: 0.3925  
- acc: 0.9003 - val\_loss: 0.1973 - val\_acc: 0.9544

Epoch 13/40

60000/60000 [=====] - 6s 99us/step - loss: 0.3740  
- acc: 0.9039 - val\_loss: 0.1941 - val\_acc: 0.9564

Epoch 14/40

60000/60000 [=====] - 6s 98us/step - loss: 0.3547  
- acc: 0.9098 - val\_loss: 0.1894 - val\_acc: 0.9559

Epoch 15/40

60000/60000 [=====] - 6s 99us/step - loss: 0.3401  
- acc: 0.9157 - val\_loss: 0.1811 - val\_acc: 0.9591

Epoch 16/40

60000/60000 [=====] - 6s 98us/step - loss: 0.3280  
- acc: 0.9181 - val\_loss: 0.1781 - val\_acc: 0.9592

Epoch 17/40

60000/60000 [=====] - 6s 98us/step - loss: 0.3098  
- acc: 0.9238 - val\_loss: 0.1679 - val\_acc: 0.9614

Epoch 18/40

60000/60000 [=====] - 6s 99us/step - loss: 0.3059  
- acc: 0.9255 - val\_loss: 0.1616 - val\_acc: 0.9638

Epoch 19/40

60000/60000 [=====] - 6s 99us/step - loss: 0.2924  
- acc: 0.9280 - val\_loss: 0.1594 - val\_acc: 0.9635

Epoch 20/40

60000/60000 [=====] - 6s 100us/step - loss: 0.281  
4 - acc: 0.9311 - val\_loss: 0.1583 - val\_acc: 0.9642

Epoch 21/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.2781  
- acc: 0.9325 - val\_loss: 0.1605 - val\_acc: 0.9650

Epoch 22/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.2714  
- acc: 0.9347 - val\_loss: 0.1472 - val\_acc: 0.9663

Epoch 23/40  
60000/60000 [=====] - 6s 100us/step - loss: 0.2608  
- acc: 0.9365 - val\_loss: 0.1492 - val\_acc: 0.9677

Epoch 24/40  
60000/60000 [=====] - 6s 100us/step - loss: 0.2544  
- acc: 0.9382 - val\_loss: 0.1452 - val\_acc: 0.9686

Epoch 25/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.2501  
- acc: 0.9405 - val\_loss: 0.1416 - val\_acc: 0.9688

Epoch 26/40  
60000/60000 [=====] - 6s 98us/step - loss: 0.2453  
- acc: 0.9404 - val\_loss: 0.1449 - val\_acc: 0.9682

Epoch 27/40  
60000/60000 [=====] - 6s 97us/step - loss: 0.2405  
- acc: 0.9423 - val\_loss: 0.1428 - val\_acc: 0.9685

Epoch 28/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.2320  
- acc: 0.9443 - val\_loss: 0.1358 - val\_acc: 0.9705

Epoch 29/40  
60000/60000 [=====] - 6s 98us/step - loss: 0.2278  
- acc: 0.9451 - val\_loss: 0.1319 - val\_acc: 0.9698

Epoch 30/40  
60000/60000 [=====] - 6s 97us/step - loss: 0.2251  
- acc: 0.9462 - val\_loss: 0.1301 - val\_acc: 0.9706

Epoch 31/40  
60000/60000 [=====] - 6s 97us/step - loss: 0.2198  
- acc: 0.9470 - val\_loss: 0.1358 - val\_acc: 0.9710

Epoch 32/40  
60000/60000 [=====] - 6s 97us/step - loss: 0.2136  
- acc: 0.9489 - val\_loss: 0.1295 - val\_acc: 0.9713

Epoch 33/40  
60000/60000 [=====] - 6s 97us/step - loss: 0.2114  
- acc: 0.9481 - val\_loss: 0.1252 - val\_acc: 0.9718

Epoch 34/40  
60000/60000 [=====] - 6s 98us/step - loss: 0.2074  
- acc: 0.9498 - val\_loss: 0.1268 - val\_acc: 0.9720

Epoch 35/40  
60000/60000 [=====] - 6s 98us/step - loss: 0.2042  
- acc: 0.9505 - val\_loss: 0.1261 - val\_acc: 0.9722

Epoch 36/40  
60000/60000 [=====] - 6s 98us/step - loss: 0.1955  
- acc: 0.9524 - val\_loss: 0.1285 - val\_acc: 0.9741

Epoch 37/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.1988  
- acc: 0.9526 - val\_loss: 0.1247 - val\_acc: 0.9730

Epoch 38/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.1964  
- acc: 0.9519 - val\_loss: 0.1229 - val\_acc: 0.9731

Epoch 39/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.1906  
- acc: 0.9535 - val\_loss: 0.1247 - val\_acc: 0.9728

Epoch 40/40  
60000/60000 [=====] - 6s 99us/step - loss: 0.1880  
- acc: 0.9551 - val\_loss: 0.1221 - val\_acc: 0.9721

In [38]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.12213036285452544

Test accuracy: 0.9721





## Tanh activation +adam optimizer+ Batch Normalization + Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2(128 neurons))---(hidden layer 3 (256 neurons))---(output layer)]

In [26]:

```
model_final = Sequential()
model_final.add(Dense(512, activation='tanh', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_final.add(BatchNormalization())

model_final.add(Dense(128, activation='tanh', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())

model_final.add(Dense(256, activation='tanh', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_final.add(BatchNormalization())

model_final.add(Dense(output_dim, activation='softmax'))
model_final.summary()
```

W1003 04:23:29.100877 14324 deprecation\_wrapper.py:119] From C:\Users\RASHU TYAGI\Anaconda3\lib\site-packages\keras\backend\tensorflow\_backend.py:133: The name tf.placeholder\_with\_default is deprecated. Please use tf.compat.v1.placeholder\_with\_default instead.

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_4 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_5 (Dense)	(None, 256)	33024
batch_normalization_3 (Batch Normalization)	(None, 256)	1024
dense_6 (Dense)	(None, 10)	2570
Total params: 506,762		
Trainable params: 504,970		
Non-trainable params: 1,792		

In [27]:

```
nb_epoch = 30
batch_size = 64
```

In [28]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model_final.fit(X_train, Y_train, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
W1003 04:23:33.710057 14324 deprecation.py:323] From C:\Users\RASHU TYAGI
\Anaconda3\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_
dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) i
s deprecated and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/30

60000/60000 [=====] - 13s 216us/step - loss: 0.43

63 - acc: 0.8650 - val\_loss: 0.2711 - val\_acc: 0.9171

Epoch 2/30

60000/60000 [=====] - 11s 182us/step - loss: 0.27

92 - acc: 0.9136 - val\_loss: 0.2183 - val\_acc: 0.9335

Epoch 3/30

60000/60000 [=====] - 11s 183us/step - loss: 0.23

51 - acc: 0.9277 - val\_loss: 0.1845 - val\_acc: 0.9446

Epoch 4/30

60000/60000 [=====] - 11s 181us/step - loss: 0.20

73 - acc: 0.9361 - val\_loss: 0.1667 - val\_acc: 0.9490

Epoch 5/30

60000/60000 [=====] - 11s 181us/step - loss: 0.18

65 - acc: 0.9439 - val\_loss: 0.1544 - val\_acc: 0.9537

Epoch 6/30

60000/60000 [=====] - 11s 181us/step - loss: 0.16

82 - acc: 0.9478 - val\_loss: 0.1342 - val\_acc: 0.9585

Epoch 7/30

60000/60000 [=====] - 11s 182us/step - loss: 0.15

84 - acc: 0.9523 - val\_loss: 0.1270 - val\_acc: 0.9631

Epoch 8/30

60000/60000 [=====] - 11s 183us/step - loss: 0.14

99 - acc: 0.9532 - val\_loss: 0.1319 - val\_acc: 0.9611

Epoch 9/30

60000/60000 [=====] - 11s 181us/step - loss: 0.14

27 - acc: 0.9554 - val\_loss: 0.1290 - val\_acc: 0.9593

Epoch 10/30

60000/60000 [=====] - 11s 183us/step - loss: 0.13

42 - acc: 0.9582 - val\_loss: 0.1193 - val\_acc: 0.9643

Epoch 11/30

60000/60000 [=====] - 11s 182us/step - loss: 0.12

36 - acc: 0.9606 - val\_loss: 0.1210 - val\_acc: 0.9629

Epoch 12/30

60000/60000 [=====] - 11s 181us/step - loss: 0.12

00 - acc: 0.9621 - val\_loss: 0.1236 - val\_acc: 0.9624

Epoch 13/30

60000/60000 [=====] - 11s 180us/step - loss: 0.11

60 - acc: 0.9641 - val\_loss: 0.1142 - val\_acc: 0.9643

Epoch 14/30

60000/60000 [=====] - 11s 183us/step - loss: 0.11

35 - acc: 0.9649 - val\_loss: 0.1136 - val\_acc: 0.9672

Epoch 15/30

60000/60000 [=====] - 11s 180us/step - loss: 0.10

95 - acc: 0.9655 - val\_loss: 0.1187 - val\_acc: 0.9642

Epoch 16/30

60000/60000 [=====] - 11s 181us/step - loss: 0.10

36 - acc: 0.9672 - val\_loss: 0.1100 - val\_acc: 0.9653

Epoch 17/30

60000/60000 [=====] - 11s 182us/step - loss: 0.09

68 - acc: 0.9700 - val\_loss: 0.1026 - val\_acc: 0.9673

Epoch 18/30

60000/60000 [=====] - 11s 183us/step - loss: 0.09

74 - acc: 0.9692 - val\_loss: 0.0963 - val\_acc: 0.9703

Epoch 19/30

60000/60000 [=====] - 11s 183us/step - loss: 0.09

40 - acc: 0.9707 - val\_loss: 0.1024 - val\_acc: 0.9702

Epoch 20/30

60000/60000 [=====] - 11s 183us/step - loss: 0.09

14 - acc: 0.9698 - val\_loss: 0.0979 - val\_acc: 0.9720

Epoch 21/30  
60000/60000 [=====] - 11s 179us/step - loss: 0.08  
58 - acc: 0.9730 - val\_loss: 0.1003 - val\_acc: 0.9719

Epoch 22/30  
60000/60000 [=====] - 11s 182us/step - loss: 0.08  
63 - acc: 0.9726 - val\_loss: 0.0937 - val\_acc: 0.9727

Epoch 23/30  
60000/60000 [=====] - 11s 182us/step - loss: 0.08  
28 - acc: 0.9742 - val\_loss: 0.1041 - val\_acc: 0.9678

Epoch 24/30  
60000/60000 [=====] - 11s 179us/step - loss: 0.07  
90 - acc: 0.9749 - val\_loss: 0.0926 - val\_acc: 0.9730

Epoch 25/30  
60000/60000 [=====] - 11s 180us/step - loss: 0.07  
56 - acc: 0.9764 - val\_loss: 0.0916 - val\_acc: 0.9719

Epoch 26/30  
60000/60000 [=====] - 11s 181us/step - loss: 0.07  
38 - acc: 0.9760 - val\_loss: 0.0939 - val\_acc: 0.9721

Epoch 27/30  
60000/60000 [=====] - 11s 183us/step - loss: 0.07  
11 - acc: 0.9779 - val\_loss: 0.0980 - val\_acc: 0.9716

Epoch 28/30  
60000/60000 [=====] - 11s 182us/step - loss: 0.07  
16 - acc: 0.9771 - val\_loss: 0.0909 - val\_acc: 0.9732

Epoch 29/30  
60000/60000 [=====] - 11s 181us/step - loss: 0.07  
13 - acc: 0.9773 - val\_loss: 0.0945 - val\_acc: 0.9720

Epoch 30/30  
60000/60000 [=====] - 11s 180us/step - loss: 0.06  
75 - acc: 0.9786 - val\_loss: 0.0896 - val\_acc: 0.9736

In [29]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08960605073939078

Test accuracy: 0.9736



## Next architecture

**glorot\_uniform initialization + sgd optimizer + relu activation function + Dropout(0.33) + no batch normalization + Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))--(hidden layer 4 (1024 neurons))----(hidden layer 5 (2048 neurons))---(output layer)]**

In [36]:

```
import keras.utils

model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer = keras.initializers.glorot_normal(seed=None)))
model_final.add(Dropout(0.33))

model_final.add(Dense(128, activation='relu', kernel_initializer = keras.initializers.glorot_normal(seed=None)))
model_final.add(Dropout(0.33))

model_final.add(Dense(256, activation='relu', kernel_initializer = keras.initializers.glorot_normal(seed=None)))
model_final.add(Dropout(0.33))

model_final.add(Dense(1024, activation='relu', kernel_initializer = keras.initializers.glorot_normal(seed=None)))
model_final.add(Dropout(0.33))

model_final.add(Dense(2048, activation='relu', kernel_initializer = keras.initializers.glorot_normal(seed=None)))
model_final.add(Dropout(0.33))

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dropout_2 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 128)	65664
dropout_3 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 256)	33024
dropout_4 (Dropout)	(None, 256)	0
dense_11 (Dense)	(None, 1024)	263168
dropout_5 (Dropout)	(None, 1024)	0
dense_12 (Dense)	(None, 2048)	2099200
dropout_6 (Dropout)	(None, 2048)	0
dense_13 (Dense)	(None, 10)	20490
Total params: 2,883,466		
Trainable params: 2,883,466		
Non-trainable params: 0		



In [37]:

```
nb_epoch = 60
```

In [38]:

```
model_final.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/60

60000/60000 [=====] - 4s 68us/step - loss: 1.8581  
- acc: 0.3439 - val\_loss: 0.7089 - val\_acc: 0.7973

Epoch 2/60

60000/60000 [=====] - 3s 58us/step - loss: 0.7309  
- acc: 0.7606 - val\_loss: 0.3608 - val\_acc: 0.8915

Epoch 3/60

60000/60000 [=====] - 4s 60us/step - loss: 0.4883  
- acc: 0.8503 - val\_loss: 0.2734 - val\_acc: 0.9179

Epoch 4/60

60000/60000 [=====] - 3s 58us/step - loss: 0.3872  
- acc: 0.8858 - val\_loss: 0.2224 - val\_acc: 0.9335

Epoch 5/60

60000/60000 [=====] - 3s 58us/step - loss: 0.3271  
- acc: 0.9028 - val\_loss: 0.1916 - val\_acc: 0.9418

Epoch 6/60

60000/60000 [=====] - 3s 58us/step - loss: 0.2824  
- acc: 0.9163 - val\_loss: 0.1719 - val\_acc: 0.9490

Epoch 7/60

60000/60000 [=====] - 3s 58us/step - loss: 0.2523  
- acc: 0.9250 - val\_loss: 0.1520 - val\_acc: 0.9535

Epoch 8/60

60000/60000 [=====] - 3s 58us/step - loss: 0.2295  
- acc: 0.9324 - val\_loss: 0.1409 - val\_acc: 0.9566

Epoch 9/60

60000/60000 [=====] - 3s 58us/step - loss: 0.2080  
- acc: 0.9387 - val\_loss: 0.1292 - val\_acc: 0.9607

Epoch 10/60

60000/60000 [=====] - 4s 60us/step - loss: 0.1938  
- acc: 0.9428 - val\_loss: 0.1225 - val\_acc: 0.9638

Epoch 11/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1786  
- acc: 0.9467 - val\_loss: 0.1157 - val\_acc: 0.9645

Epoch 12/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1687  
- acc: 0.9501 - val\_loss: 0.1106 - val\_acc: 0.9675

Epoch 13/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1554  
- acc: 0.9542 - val\_loss: 0.1050 - val\_acc: 0.9681

Epoch 14/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1510  
- acc: 0.9558 - val\_loss: 0.0996 - val\_acc: 0.9708

Epoch 15/60

60000/60000 [=====] - 3s 58us/step - loss: 0.1392  
- acc: 0.9590 - val\_loss: 0.0981 - val\_acc: 0.9707

Epoch 16/60

60000/60000 [=====] - 4s 59us/step - loss: 0.1329  
- acc: 0.9613 - val\_loss: 0.0967 - val\_acc: 0.9708

Epoch 17/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1241  
- acc: 0.9637 - val\_loss: 0.0930 - val\_acc: 0.9725

Epoch 18/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1197  
- acc: 0.9646 - val\_loss: 0.0902 - val\_acc: 0.9734

Epoch 19/60

60000/60000 [=====] - 3s 57us/step - loss: 0.1135  
- acc: 0.9668 - val\_loss: 0.0868 - val\_acc: 0.9748

Epoch 20/60

60000/60000 [=====] - 4s 60us/step - loss: 0.1095  
- acc: 0.9673 - val\_loss: 0.0877 - val\_acc: 0.9745

Epoch 21/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.1042  
- acc: 0.9685 - val\_loss: 0.0897 - val\_acc: 0.9737

Epoch 22/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.1006  
- acc: 0.9704 - val\_loss: 0.0831 - val\_acc: 0.9761

Epoch 23/60  
60000/60000 [=====] - 4s 58us/step - loss: 0.0971  
- acc: 0.9711 - val\_loss: 0.0829 - val\_acc: 0.9767

Epoch 24/60  
60000/60000 [=====] - 4s 58us/step - loss: 0.0904  
- acc: 0.9729 - val\_loss: 0.0842 - val\_acc: 0.9760

Epoch 25/60  
60000/60000 [=====] - 4s 60us/step - loss: 0.0878  
- acc: 0.9730 - val\_loss: 0.0828 - val\_acc: 0.9770

Epoch 26/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0839  
- acc: 0.9748 - val\_loss: 0.0795 - val\_acc: 0.9779

Epoch 27/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0817  
- acc: 0.9761 - val\_loss: 0.0798 - val\_acc: 0.9780

Epoch 28/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0829  
- acc: 0.9749 - val\_loss: 0.0794 - val\_acc: 0.9783

Epoch 29/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0789  
- acc: 0.9766 - val\_loss: 0.0783 - val\_acc: 0.9796

Epoch 30/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0754  
- acc: 0.9772 - val\_loss: 0.0780 - val\_acc: 0.9788

Epoch 31/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0699  
- acc: 0.9793 - val\_loss: 0.0805 - val\_acc: 0.9782

Epoch 32/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0690  
- acc: 0.9792 - val\_loss: 0.0795 - val\_acc: 0.9785

Epoch 33/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0671  
- acc: 0.9801 - val\_loss: 0.0815 - val\_acc: 0.9776

Epoch 34/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0660  
- acc: 0.9796 - val\_loss: 0.0791 - val\_acc: 0.9802

Epoch 35/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0659  
- acc: 0.9800 - val\_loss: 0.0741 - val\_acc: 0.9800

Epoch 36/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0640  
- acc: 0.9808 - val\_loss: 0.0777 - val\_acc: 0.9798

Epoch 37/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0603  
- acc: 0.9818 - val\_loss: 0.0768 - val\_acc: 0.9801

Epoch 38/60  
60000/60000 [=====] - 4s 59us/step - loss: 0.0594  
- acc: 0.9817 - val\_loss: 0.0752 - val\_acc: 0.9812

Epoch 39/60  
60000/60000 [=====] - 4s 60us/step - loss: 0.0591  
- acc: 0.9822 - val\_loss: 0.0752 - val\_acc: 0.9801

Epoch 40/60  
60000/60000 [=====] - 4s 63us/step - loss: 0.0543  
- acc: 0.9835 - val\_loss: 0.0750 - val\_acc: 0.9806

Epoch 41/60

```
60000/60000 [=====] - 4s 59us/step - loss: 0.0518
- acc: 0.9849 - val_loss: 0.0762 - val_acc: 0.9808
Epoch 42/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0529
- acc: 0.9832 - val_loss: 0.0789 - val_acc: 0.9795
Epoch 43/60
60000/60000 [=====] - 4s 60us/step - loss: 0.0500
- acc: 0.9843 - val_loss: 0.0788 - val_acc: 0.9804
Epoch 44/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0498
- acc: 0.9844 - val_loss: 0.0746 - val_acc: 0.9818
Epoch 45/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0497
- acc: 0.9844 - val_loss: 0.0731 - val_acc: 0.9817
Epoch 46/60
60000/60000 [=====] - 4s 58us/step - loss: 0.0499
- acc: 0.9841 - val_loss: 0.0773 - val_acc: 0.9808
Epoch 47/60
60000/60000 [=====] - 4s 58us/step - loss: 0.0450
- acc: 0.9860 - val_loss: 0.0758 - val_acc: 0.9816
Epoch 48/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0450
- acc: 0.9856 - val_loss: 0.0760 - val_acc: 0.9811
Epoch 49/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0445
- acc: 0.9859 - val_loss: 0.0747 - val_acc: 0.9824
Epoch 50/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0427
- acc: 0.9868 - val_loss: 0.0748 - val_acc: 0.9815
Epoch 51/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0409
- acc: 0.9875 - val_loss: 0.0769 - val_acc: 0.9826
Epoch 52/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0415
- acc: 0.9870 - val_loss: 0.0771 - val_acc: 0.9823
Epoch 53/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0386
- acc: 0.9881 - val_loss: 0.0758 - val_acc: 0.9825
Epoch 54/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0379
- acc: 0.9881 - val_loss: 0.0783 - val_acc: 0.9812
Epoch 55/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0387
- acc: 0.9878 - val_loss: 0.0735 - val_acc: 0.9828
Epoch 56/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0353
- acc: 0.9890 - val_loss: 0.0728 - val_acc: 0.9826
Epoch 57/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0373
- acc: 0.9884 - val_loss: 0.0771 - val_acc: 0.9826
Epoch 58/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0358
- acc: 0.9891 - val_loss: 0.0754 - val_acc: 0.9828
Epoch 59/60
60000/60000 [=====] - 4s 59us/step - loss: 0.0361
- acc: 0.9890 - val_loss: 0.0749 - val_acc: 0.9832
Epoch 60/60
60000/60000 [=====] - 3s 58us/step - loss: 0.0341
- acc: 0.9892 - val_loss: 0.0759 - val_acc: 0.9832
```

In [39]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07590167486594146

Test accuracy: 0.9832



## Next Architecture

**he\_normal initialization + adam optimizer + softmax activation function + Dropout(0.66) + batch normalization(to only 2 layers) + Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))--(hidden layer 4 (1024 neurons))----(hidden layer 5 (2048 neurons))---(output layer)]**

In [40]:

```
# refer this for he_normal initialization --> https://keras.io/activations/
import keras.utils

model_final = Sequential()

model_final.add(Dense(512, activation='softmax', input_shape=(input_dim,), kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))

model_final.add(Dense(128, activation='softmax', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))
model_final.add(BatchNormalization())

model_final.add(Dense(256, activation='softmax', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))

model_final.add(Dense(1024, activation='softmax', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))

model_final.add(Dense(2048, activation='softmax', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))
model_final.add(BatchNormalization())

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()
```



```

W1003 06:40:12.831539 14324 nn_ops.py:4224] Large dropout rate: 0.66 (>0.
5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. P
lease ensure that this is intended.
W1003 06:40:12.853988 14324 nn_ops.py:4224] Large dropout rate: 0.66 (>0.
5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. P
lease ensure that this is intended.
W1003 06:40:12.924259 14324 nn_ops.py:4224] Large dropout rate: 0.66 (>0.
5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. P
lease ensure that this is intended.
W1003 06:40:12.944783 14324 nn_ops.py:4224] Large dropout rate: 0.66 (>0.
5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. P
lease ensure that this is intended.
W1003 06:40:12.964304 14324 nn_ops.py:4224] Large dropout rate: 0.66 (>0.
5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. P
lease ensure that this is intended.

```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
dropout_7 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 128)	65664
dropout_8 (Dropout)	(None, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 256)	33024
dropout_9 (Dropout)	(None, 256)	0
dense_17 (Dense)	(None, 1024)	263168
dropout_10 (Dropout)	(None, 1024)	0
dense_18 (Dense)	(None, 2048)	2099200
dropout_11 (Dropout)	(None, 2048)	0
batch_normalization_5 (Batch Normalization)	(None, 2048)	8192
dense_19 (Dense)	(None, 10)	20490
Total params: 2,892,170		
Trainable params: 2,887,818		
Non-trainable params: 4,352		

In [42]:

```

nb_epoch = 50
batch_size = 256

```

In [43]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
  
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/50

60000/60000 [=====] - 3s 56us/step - loss: 2.3027  
- acc: 0.1090 - val\_loss: 2.3015 - val\_acc: 0.1028

Epoch 2/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1086 - val\_loss: 2.3016 - val\_acc: 0.1135

Epoch 3/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1083 - val\_loss: 2.3014 - val\_acc: 0.1135

Epoch 4/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3028  
- acc: 0.1077 - val\_loss: 2.3014 - val\_acc: 0.1135

Epoch 5/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3027  
- acc: 0.1073 - val\_loss: 2.3025 - val\_acc: 0.1135

Epoch 6/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3028  
- acc: 0.1098 - val\_loss: 2.3018 - val\_acc: 0.1028

Epoch 7/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3030  
- acc: 0.1076 - val\_loss: 2.3031 - val\_acc: 0.1135

Epoch 8/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3031  
- acc: 0.1094 - val\_loss: 2.3036 - val\_acc: 0.1135

Epoch 9/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1077 - val\_loss: 2.3021 - val\_acc: 0.1135

Epoch 10/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1077 - val\_loss: 2.3017 - val\_acc: 0.1135

Epoch 11/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1080 - val\_loss: 2.3014 - val\_acc: 0.1135

Epoch 12/50

60000/60000 [=====] - 2s 39us/step - loss: 2.3031  
- acc: 0.1073 - val\_loss: 2.3030 - val\_acc: 0.1028

Epoch 13/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3031  
- acc: 0.1081 - val\_loss: 2.3013 - val\_acc: 0.1135

Epoch 14/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3034  
- acc: 0.1084 - val\_loss: 2.3015 - val\_acc: 0.1135

Epoch 15/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1073 - val\_loss: 2.3024 - val\_acc: 0.1135

Epoch 16/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3036  
- acc: 0.1070 - val\_loss: 2.3023 - val\_acc: 0.1135

Epoch 17/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1077 - val\_loss: 2.3015 - val\_acc: 0.1009

Epoch 18/50

60000/60000 [=====] - 2s 38us/step - loss: 2.3033  
- acc: 0.1062 - val\_loss: 2.3027 - val\_acc: 0.1135

Epoch 19/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3035  
- acc: 0.1063 - val\_loss: 2.3025 - val\_acc: 0.1135

Epoch 20/50

60000/60000 [=====] - 2s 37us/step - loss: 2.3033  
- acc: 0.1075 - val\_loss: 2.3014 - val\_acc: 0.1135

Epoch 21/50  
60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1067 - val\_loss: 2.3020 - val\_acc: 0.1135

Epoch 22/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3029  
- acc: 0.1083 - val\_loss: 2.3026 - val\_acc: 0.1135

Epoch 23/50  
60000/60000 [=====] - 2s 38us/step - loss: 2.3033  
- acc: 0.1078 - val\_loss: 2.3022 - val\_acc: 0.1010

Epoch 24/50  
60000/60000 [=====] - 2s 38us/step - loss: 2.3029  
- acc: 0.1087 - val\_loss: 2.3014 - val\_acc: 0.1135

Epoch 25/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3030  
- acc: 0.1084 - val\_loss: 2.3027 - val\_acc: 0.1135

Epoch 26/50  
60000/60000 [=====] - 2s 38us/step - loss: 2.3032  
- acc: 0.1071 - val\_loss: 2.3026 - val\_acc: 0.1032

Epoch 27/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3035  
- acc: 0.1079 - val\_loss: 2.3020 - val\_acc: 0.1010

Epoch 28/50  
60000/60000 [=====] - 2s 38us/step - loss: 2.3031  
- acc: 0.1080 - val\_loss: 2.3016 - val\_acc: 0.1135

Epoch 29/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1073 - val\_loss: 2.3024 - val\_acc: 0.1135

Epoch 30/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1086 - val\_loss: 2.3016 - val\_acc: 0.1135

Epoch 31/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3030  
- acc: 0.1077 - val\_loss: 2.3016 - val\_acc: 0.1135

Epoch 32/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3031  
- acc: 0.1077 - val\_loss: 2.3020 - val\_acc: 0.1135

Epoch 33/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1087 - val\_loss: 2.3024 - val\_acc: 0.1135

Epoch 34/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1076 - val\_loss: 2.3022 - val\_acc: 0.1135

Epoch 35/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3029  
- acc: 0.1081 - val\_loss: 2.3022 - val\_acc: 0.1135

Epoch 36/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3031  
- acc: 0.1089 - val\_loss: 2.3018 - val\_acc: 0.1135

Epoch 37/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3032  
- acc: 0.1080 - val\_loss: 2.3021 - val\_acc: 0.1028

Epoch 38/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3031  
- acc: 0.1070 - val\_loss: 2.3022 - val\_acc: 0.1028

Epoch 39/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3028  
- acc: 0.1066 - val\_loss: 2.3026 - val\_acc: 0.1135

Epoch 40/50  
60000/60000 [=====] - 2s 37us/step - loss: 2.3030  
- acc: 0.1068 - val\_loss: 2.3022 - val\_acc: 0.1135

Epoch 41/50

```
60000/60000 [=====] - 2s 37us/step - loss: 2.3030
- acc: 0.1075 - val_loss: 2.3018 - val_acc: 0.1135
Epoch 42/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3033
- acc: 0.1065 - val_loss: 2.3017 - val_acc: 0.1135
Epoch 43/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3031
- acc: 0.1091 - val_loss: 2.3024 - val_acc: 0.1028
Epoch 44/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3033
- acc: 0.1067 - val_loss: 2.3015 - val_acc: 0.1135
Epoch 45/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3027
- acc: 0.1081 - val_loss: 2.3035 - val_acc: 0.1135
Epoch 46/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3032
- acc: 0.1068 - val_loss: 2.3025 - val_acc: 0.1135
Epoch 47/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3029
- acc: 0.1072 - val_loss: 2.3015 - val_acc: 0.1135
Epoch 48/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3028
- acc: 0.1091 - val_loss: 2.3025 - val_acc: 0.1028
Epoch 49/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3032
- acc: 0.1079 - val_loss: 2.3015 - val_acc: 0.1135
Epoch 50/50
60000/60000 [=====] - 2s 37us/step - loss: 2.3030
- acc: 0.1079 - val_loss: 2.3017 - val_acc: 0.1010
```

In [44]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))

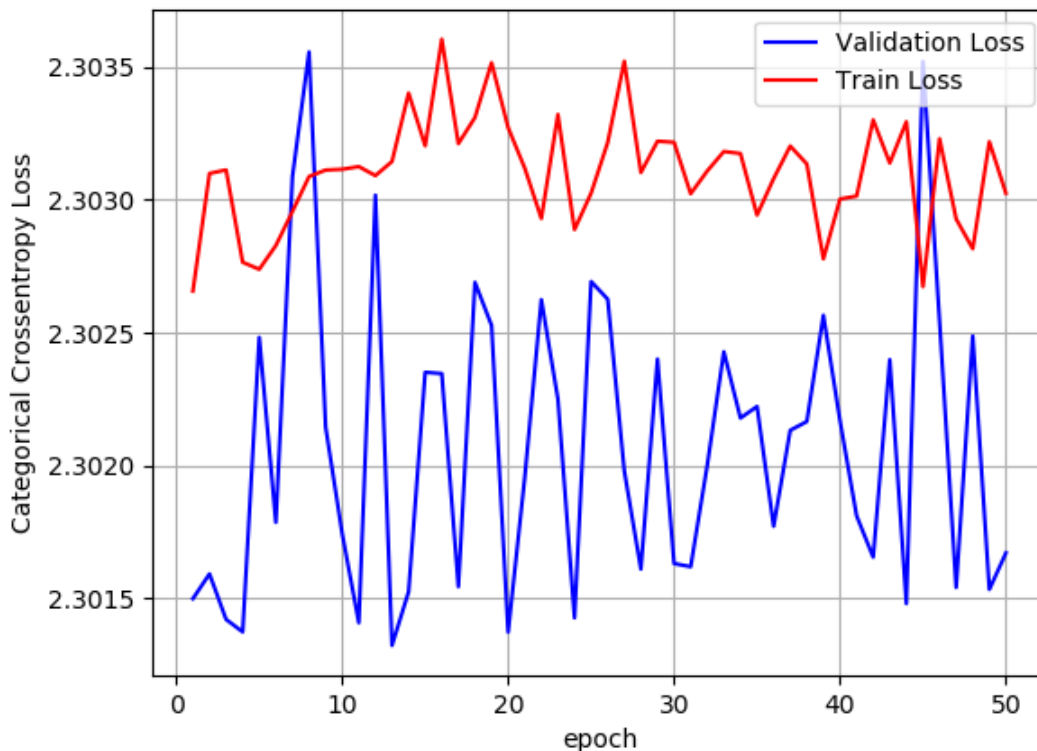
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 2.3016698822021486

Test accuracy: 0.101



**We received a very poor accuracy over there let us try the same code with relu activation**

In [46]:

```
# refer this for he_normal initialization --> https://keras.io/activations/
import keras.utils

model_final = Sequential()

model_final.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))

model_final.add(Dense(128, activation='relu', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))
model_final.add(BatchNormalization())

model_final.add(Dense(256, activation='relu', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))

model_final.add(Dense(1024, activation='relu', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))

model_final.add(Dense(2048, activation='relu', kernel_initializer = keras.initializers.he_normal(seed=None)))
model_final.add(Dropout(0.66))
model_final.add(BatchNormalization())

model_final.add(Dense(output_dim, activation='softmax'))

model_final.summary()
```



Layer (type)	Output Shape	Param #
=====	=====	=====
dense_20 (Dense)	(None, 512)	401920
dropout_12 (Dropout)	(None, 512)	0
dense_21 (Dense)	(None, 128)	65664
dropout_13 (Dropout)	(None, 128)	0
batch_normalization_6 (Batch Normalization)	(None, 128)	512
dense_22 (Dense)	(None, 256)	33024
dropout_14 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 1024)	263168
dropout_15 (Dropout)	(None, 1024)	0
dense_24 (Dense)	(None, 2048)	2099200
dropout_16 (Dropout)	(None, 2048)	0
batch_normalization_7 (Batch Normalization)	(None, 2048)	8192
dense_25 (Dense)	(None, 10)	20490
=====	=====	=====
Total params: 2,892,170		
Trainable params: 2,887,818		
Non-trainable params: 4,352		

In [47]:

```
nb_epoch = 50
batch_size = 256
```

In [48]:

```
model_final.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_final.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/50

60000/60000 [=====] - 3s 55us/step - loss: 1.9138  
- acc: 0.3571 - val\_loss: 1.3399 - val\_acc: 0.4876

Epoch 2/50

60000/60000 [=====] - 2s 35us/step - loss: 0.9259  
- acc: 0.6837 - val\_loss: 0.5729 - val\_acc: 0.7837

Epoch 3/50

60000/60000 [=====] - 2s 35us/step - loss: 0.6319  
- acc: 0.8045 - val\_loss: 0.4046 - val\_acc: 0.8502

Epoch 4/50

60000/60000 [=====] - 2s 35us/step - loss: 0.5001  
- acc: 0.8528 - val\_loss: 0.2862 - val\_acc: 0.9020

Epoch 5/50

60000/60000 [=====] - 2s 35us/step - loss: 0.4270  
- acc: 0.8785 - val\_loss: 0.2521 - val\_acc: 0.9145

Epoch 6/50

60000/60000 [=====] - 2s 35us/step - loss: 0.3747  
- acc: 0.8951 - val\_loss: 0.1859 - val\_acc: 0.9466

Epoch 7/50

60000/60000 [=====] - 2s 36us/step - loss: 0.3425  
- acc: 0.9070 - val\_loss: 0.1850 - val\_acc: 0.9443

Epoch 8/50

60000/60000 [=====] - 2s 36us/step - loss: 0.3117  
- acc: 0.9154 - val\_loss: 0.1650 - val\_acc: 0.9538

Epoch 9/50

60000/60000 [=====] - 2s 36us/step - loss: 0.2962  
- acc: 0.9197 - val\_loss: 0.1580 - val\_acc: 0.9577

Epoch 10/50

60000/60000 [=====] - 2s 36us/step - loss: 0.2744  
- acc: 0.9270 - val\_loss: 0.1428 - val\_acc: 0.9624

Epoch 11/50

60000/60000 [=====] - 2s 37us/step - loss: 0.2594  
- acc: 0.9315 - val\_loss: 0.1284 - val\_acc: 0.9661

Epoch 12/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2456  
- acc: 0.9355 - val\_loss: 0.1291 - val\_acc: 0.9655

Epoch 13/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2358  
- acc: 0.9380 - val\_loss: 0.1303 - val\_acc: 0.9679

Epoch 14/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2246  
- acc: 0.9410 - val\_loss: 0.1205 - val\_acc: 0.9693

Epoch 15/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2226  
- acc: 0.9420 - val\_loss: 0.1140 - val\_acc: 0.9704

Epoch 16/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2121  
- acc: 0.9449 - val\_loss: 0.1218 - val\_acc: 0.9692

Epoch 17/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2031  
- acc: 0.9496 - val\_loss: 0.1091 - val\_acc: 0.9726

Epoch 18/50

60000/60000 [=====] - 2s 35us/step - loss: 0.2023  
- acc: 0.9479 - val\_loss: 0.1066 - val\_acc: 0.9731

Epoch 19/50

60000/60000 [=====] - 2s 35us/step - loss: 0.1920  
- acc: 0.9505 - val\_loss: 0.1035 - val\_acc: 0.9730

Epoch 20/50

60000/60000 [=====] - 2s 35us/step - loss: 0.1936  
- acc: 0.9506 - val\_loss: 0.1092 - val\_acc: 0.9728

Epoch 21/50  
60000/60000 [=====] - 2s 36us/step - loss: 0.1823  
- acc: 0.9532 - val\_loss: 0.1085 - val\_acc: 0.9733

Epoch 22/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1784  
- acc: 0.9547 - val\_loss: 0.1020 - val\_acc: 0.9730

Epoch 23/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1742  
- acc: 0.9534 - val\_loss: 0.0984 - val\_acc: 0.9756

Epoch 24/50  
60000/60000 [=====] - 2s 38us/step - loss: 0.1691  
- acc: 0.9572 - val\_loss: 0.1022 - val\_acc: 0.9739

Epoch 25/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1717  
- acc: 0.9562 - val\_loss: 0.0985 - val\_acc: 0.9740

Epoch 26/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1632  
- acc: 0.9579 - val\_loss: 0.1000 - val\_acc: 0.9758

Epoch 27/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1612  
- acc: 0.9577 - val\_loss: 0.0976 - val\_acc: 0.9758

Epoch 28/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1507  
- acc: 0.9603 - val\_loss: 0.0916 - val\_acc: 0.9773

Epoch 29/50  
60000/60000 [=====] - 2s 36us/step - loss: 0.1564  
- acc: 0.9597 - val\_loss: 0.0921 - val\_acc: 0.9776

Epoch 30/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1506  
- acc: 0.9613 - val\_loss: 0.0941 - val\_acc: 0.9764

Epoch 31/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1488  
- acc: 0.9619 - val\_loss: 0.0884 - val\_acc: 0.9792

Epoch 32/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1476  
- acc: 0.9614 - val\_loss: 0.0891 - val\_acc: 0.9775

Epoch 33/50  
60000/60000 [=====] - 2s 36us/step - loss: 0.1408  
- acc: 0.9639 - val\_loss: 0.0893 - val\_acc: 0.9778

Epoch 34/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1410  
- acc: 0.9637 - val\_loss: 0.0886 - val\_acc: 0.9779

Epoch 35/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1394  
- acc: 0.9649 - val\_loss: 0.0895 - val\_acc: 0.9778

Epoch 36/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1389  
- acc: 0.9648 - val\_loss: 0.0909 - val\_acc: 0.9772

Epoch 37/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1329  
- acc: 0.9658 - val\_loss: 0.0883 - val\_acc: 0.9782

Epoch 38/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1335  
- acc: 0.9660 - val\_loss: 0.0856 - val\_acc: 0.9784

Epoch 39/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1273  
- acc: 0.9676 - val\_loss: 0.0894 - val\_acc: 0.9790

Epoch 40/50  
60000/60000 [=====] - 2s 35us/step - loss: 0.1244  
- acc: 0.9679 - val\_loss: 0.0860 - val\_acc: 0.9792

Epoch 41/50

```
60000/60000 [=====] - 2s 35us/step - loss: 0.1275
- acc: 0.9684 - val_loss: 0.0930 - val_acc: 0.9778
Epoch 42/50
60000/60000 [=====] - 2s 35us/step - loss: 0.1289
- acc: 0.9673 - val_loss: 0.0917 - val_acc: 0.9783
Epoch 43/50
60000/60000 [=====] - 2s 35us/step - loss: 0.1200
- acc: 0.9695 - val_loss: 0.0849 - val_acc: 0.9805
Epoch 44/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1201
- acc: 0.9697 - val_loss: 0.0858 - val_acc: 0.9791
Epoch 45/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1190
- acc: 0.9702 - val_loss: 0.0848 - val_acc: 0.9791
Epoch 46/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1211
- acc: 0.9691 - val_loss: 0.0846 - val_acc: 0.9795
Epoch 47/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1182
- acc: 0.9695 - val_loss: 0.0860 - val_acc: 0.9801
Epoch 48/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1164
- acc: 0.9708 - val_loss: 0.0798 - val_acc: 0.9802
Epoch 49/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1163
- acc: 0.9701 - val_loss: 0.0843 - val_acc: 0.9795
Epoch 50/50
60000/60000 [=====] - 2s 36us/step - loss: 0.1175
- acc: 0.9702 - val_loss: 0.0817 - val_acc: 0.9802
```

In [49]:

```
score = model_final.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, ve
rbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep
ochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08166552726062946

Test accuracy: 0.9802



In [53]:

```
# http://zetcode.com/python/prettytable/

from prettytable import PrettyTable

#If you get a ModuleNotFoundError error , install prettytable using: pip3 install prett
ytable

x = PrettyTable()

x.field_names = ["Architecture used", "Hidden layers", "Activation", "Optimizer used", "e
pochs", "Dropouts", "BN(batch size)", " accuracy"]

x.add_row(["784-512-128-10", 2, "RELU", "Adam" , 20, "yes", "yes(128)", 0.9825])
x.add_row(["784-512-128-256-10", 3, "RELU", "Adam" , 40, "yes", "yes(128)", 0.9833])
x.add_row(["784-512-128-256-1024-2048-10", 5, "RELU", "Adam" , 60, "yes", "yes(256)", 0.9841
])
x.add_row(["784-512-128-256-1024-2048-10", 5, "RELU", "Adam" , 40, "no", "no", 0.0958])
x.add_row(["784-512-128-256-1024-2048-10", 5, "RELU", "SGD" , 40, "no", "yes(128)", 0.9729])
x.add_row(["784-512-128-256-10", 3, "sigmoid", "Adam" , 40, "yes", "no", 0.9721])
x.add_row(["784-512-128-256-10", 3, "tanh", "Adam" , 30, "no", "yes(64)", 0.9736])
x.add_row(["784-512-128-256-1024-2048-10", 5, "RELU", "SGD" , 60, "yes", "no", 0.9832])
x.add_row(["784-512-128-256-1024-2048-10", 5, "softmax", "Adam" , 50, "yes", "yes(256)", 0.1
01])
x.add_row(["784-512-128-256-1024-2048-10", 5, "relu", "Adam" , 50, "yes", "yes(256)", 0.9802
])

print(x)
```

```
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|      Architecture used      | Hidden layers | Activation | Optimizer us
ed | epochs | Dropouts | BN(batch size) | accuracy |
+-----+-----+-----+-----+
+-----+-----+-----+-----+
|      784-512-128-10      |      2      |      RELU      |      Adam
|      20      |      yes      |      yes(128)      |      0.9825
|      784-512-128-256-10      |      3      |      RELU      |      Adam
|      40      |      yes      |      yes(128)      |      0.9833
| 784-512-128-256-1024-2048-10 |      5      |      RELU      |      Adam
|      60      |      yes      |      yes(256)      |      0.9841
| 784-512-128-256-1024-2048-10 |      5      |      RELU      |      Adam
|      40      |      no      |      no      |      0.0958
| 784-512-128-256-1024-2048-10 |      5      |      RELU      |      SGD
|      40      |      no      |      yes(128)      |      0.9729
|      784-512-128-256-10      |      3      |      sigmoid      |      Adam
|      40      |      yes      |      no      |      0.9721
|      784-512-128-256-10      |      3      |      tanh      |      Adam
|      30      |      no      |      yes(64)      |      0.9736
| 784-512-128-256-1024-2048-10 |      5      |      RELU      |      SGD
|      60      |      yes      |      no      |      0.9832
| 784-512-128-256-1024-2048-10 |      5      |      softmax      |      Adam
|      50      |      yes      |      yes(256)      |      0.101
| 784-512-128-256-1024-2048-10 |      5      |      relu      |      Adam
|      50      |      yes      |      yes(256)      |      0.9802
+-----+-----+-----+-----+
+-----+-----+-----+-----+
```

## Summary :-

1.) We tried various types of architectures for MLP with 2 hidden layers, 3 hidden layers and 5 hidden layers also.

2.) We measured the accuracy for all the cases we considered

3.) Along with different architectures we tried different initializations like xavier/cohort initialization, he\_normal initialization etc and varied them with applying batch normalization and sometimes removing batch normalizations.

4.) We also did experiments with dropout layers and we also changed the dropout rates as well.

5.) We also did experiments with the optimizers which are available to us like adam,sgd etc.

6.) Final conclusion that can be said is that choosing which will be the right architecture for us and choosing the best initialization, dropout rates, epochs, optimizers actually depends on data also and also the initialized parameters also we can see in the above experiments that some architectures gave us a very good accuracy at very low epoch also while some couldn't give us a good accuracy even at high epoch numbers hence it can be said that only increasing the epochs does not guarantee us a good accuracy after all.

7.) the graphs between the number of epochs and train/validation accuracy shows us how our model converges the loss to a value which is as minimum as possible in that epoch range.

**Best results** - If we go with kaggle results methodology which is highly dependent upon the validation accuracy considering about all our experiments although some have very close accuracy still the best accuracy was found to be 98.32% for the following configuration :-



**glorot\_uniform initialization + SGD optimizer + ReLU activation function + Dropout(0.33) + no batch normalization + Model Architecture --> [(input)--(hidden layer 1 (512 neurons))--(hidden layer 2 (128 neurons))---(hidden layer 3 (256 neurons))--(hidden layer 4 (1024 neurons))----(hidden layer 5 (2048 neurons))---(output layer)]**

In [ ]: